

# Coherence-Guided Control: Corridor Discovery and Physics-Respecting Rescue in a Bioelectric Grid

Tristan Stoltz<sup>1,\*</sup>

**1** Luminous Dynamics LLC, Dallas, TX, United States

\* tristan.stoltz@gmail.com

## Abstract

We study coherence-guided control in a bioelectric morphogenesis simulator. Soft Actor–Critic (SAC) with  $K$ -feedback discovered high-coherence corridors 50.0% of the time versus 22.2% under a reward-only baseline ( $p < 0.001$ ), a +27.8 percentage-point improvement. As a detector,  $K$  is well-discriminating and well-calibrated (ROC AUC = 0.943; PR AUC = 0.934; Brier = 0.164). In a preregistered ablation removing  $K$ -feedback entirely, discovery dropped to 12.5% (−37.5 pp vs SAC( $K$ ); +300% relative benefit with  $K$ ), confirming coherence guidance as the critical mechanism.

Under severe morphological damage, attractor-based “rescue” interventions yield lower mean performance than passive dynamics (no-rescue: 0.765 [0.738, 0.787]; rescue: 0.520 [0.360, 0.644]; Cliff’s  $\delta$  = 0.730, 95% BCa CI [0.160, 0.980], large effect), suggesting premature rescue can destabilize morphogenesis when system coherence is insufficient.

Statistical rigor is ensured via bootstrap bias-corrected accelerated (BCa) confidence intervals (10,000 stratified resamples), effect sizes (Vargha–Delaney  $A$ , Cliff’s  $\delta$ ), ROC/precision-recall analyses, and calibration assessment. All data, trained policies, analysis scripts, and SHA256 checksums are archived on OSF (DOI: [insert]) and GitHub (release v1.0-paper1).

This work establishes coherence-guided control as a principled approach for bioelectric intervention, prioritizing emergent system-level organization over local perturbations and providing reproducible design rules for navigating bioelectric state space.

## Author summary

Like planaria that regenerate heads guided by bioelectric cues, or embryos that coordinate organ formation through voltage gradients, biological systems use coherent electrical fields to navigate developmental paths. Bioelectric signals control how tissues grow and repair, but designing computational interventions is challenging: we must respect the physics of ion flows while achieving biological goals. We developed a machine learning system that navigates “corridors”—regions where bioelectric fields are coherently organized. Our key insight is to guide learning not with hand-engineered rewards, but with a coherence measure that captures collective field organization. We show that this coherence-guided approach dramatically improves corridor discovery (+300% relative to no-coherence baseline), that the coherence measure itself is an excellent predictor of success (ROC AUC = 0.943), and that attempts to “rescue” failing morphogenesis can backfire if applied when coherence is low. Statistical analyses use bootstrap confidence intervals, effect sizes, and calibration metrics to ensure rigor.

All code and data are openly available. This framework moves beyond trial-and-error perturbations toward principled bioelectric control that works *with* emergent system dynamics rather than fighting them. This suggests that AI-guided, coherence-respecting interventions—rather than aggressive forcing—could optimize regenerative outcomes while minimizing instability.

## Introduction

Bioelectric fields—generated by voltage gradients across cell membranes and gap-junctional networks—play a fundamental role in morphogenesis, regeneration, and pattern formation across phylogeny [?, ?, ?]. These fields encode positional information, coordinate cell behavior over macroscopic distances, and provide instructive cues for tissue-level organization [?, ?]. Empirical studies demonstrate causal roles for bioelectric patterns: planaria regenerate heads versus tails depending on resting potential distributions [?], and ectopic organs (eyes, limbs) can be induced by altering membrane voltage in precise spatiotemporal patterns [?]. Yet despite this foundational knowledge, the challenge of *controlling* bioelectric dynamics computationally remains formidable: perturbations must respect the underlying physics of ion flows and gap-junctional coupling [?] while steering systems toward desired morphological outcomes without destabilizing collective field organization.

To frame our approach, we define two key concepts. **Corridors** are computational constructs representing state-space regions where our simulator’s dynamics achieve high coherence ( $K > K^\dagger$ ), empirically associated with successful pattern maintenance. Experimental validation of such trajectories in real bioelectric tissues remains an open question. **Harmonies** are seven normalized components ( $H_1$ – $H_7$ ) of our coherence index  $K$ , each quantifying a distinct facet of field organization (spatial uniformity, temporal stability, boundary integrity, reciprocity, diversity, energetic cost, topological fidelity). These are combined via geometric mean to prevent compensatory trade-offs—low coherence in any dimension reduces overall  $K$ , ensuring balanced field organization.

Existing approaches to bioelectric control face a fundamental tension. On one hand, model-based methods (e.g., finite-element simulations, reaction-diffusion systems) provide mechanistic insight but require precise knowledge of network topology, ion channel kinetics, and coupling strengths—parameters often unknown or highly variable across tissues [?]. On the other hand, model-free reinforcement learning (RL) can discover policies through trial-and-error, but naive reward engineering (e.g., maximizing voltage at specific locations) frequently destabilizes collective dynamics or produces brittle solutions that fail to generalize [?]. A third option—manual trial-and-error perturbation—is labor-intensive, non-systematic, and provides limited mechanistic understanding.

Here we propose *coherence-guided control*: an RL framework that navigates bioelectric state space by optimizing for collective field *coherence* rather than hand-engineered local objectives. We hypothesize that high-coherence regions—which we term “corridors”—correspond to biophysically stable attractors that correlate with successful morphological outcomes. By constructing a geometric-mean coherence index ( $K$ ) from seven complementary measures of field organization (spatial uniformity, temporal stability, boundary integrity, etc.), we provide agents with a biophysically-grounded intrinsic reward signal. We test this framework in a 2D bioelectric grid model with realistic gap-junctional dynamics, voltage-gated ion channels, and stochastic perturbations. Our central questions are: (1) Can RL agents learn to discover high-coherence corridors more reliably than baseline heuristics? (2) Does the  $K$ -index predict morphological success with sufficient precision to serve as a control

signal? (3) What is the unique contribution of coherence feedback, isolated via ablation?  
 (4) When rescue interventions are applied to degraded morphologies, do they improve or  
 destabilize outcomes?

## Materials and Methods

### Bioelectric Grid Simulator

We simulate a 2D grid of cells with gap-junction coupling and voltage-dependent currents. At each timestep, membrane voltage evolves according to:

$$C \frac{dV_i}{dt} = I_{\text{ion}}(V_i; \theta) + \sum_{j \in \mathcal{N}(i)} g_{ij}(V_j - V_i) + I_{\text{leak}}(V_i; E_{\text{leak}}) + \eta_i(t) \quad (1)$$

where  $C$  is membrane capacitance,  $I_{\text{ion}}$  represents a nonlinear channel current term (set to zero for computational simplicity:  $\alpha = 0$ ),  $g_{ij}$  are gap-junctional conductances between 4-connected neighbors  $\mathcal{N}(i)$ ,  $I_{\text{leak}}$  enforces a biophysically bounded reversal potential  $E_{\text{leak}}$ , and  $\eta_i(t)$  is zero-mean Gaussian noise ( $\sigma = 10^{-3}$  mV,  $\sigma^2 = 10^{-6}$  mV $^2$ ). We use Neumann boundary conditions and Euler integration with  $\Delta t = 0.1$  ms.

Controllers modulate  $E_{\text{leak}}$  within physiological ranges ( $\pm 20$  mV from resting); actions are clipped to prevent nonphysical voltages. Grid dimensions: 32×32 cells; voltage range: -100 to 0 mV (hyperpolarization-only dynamics); gap-junction diffusion constant  $D$ : 0.12; leak conductance  $g$ : 0.08.

### Simulator Scope and Limitations

Our simulator employs a simplified reaction-diffusion framework with Ohmic leak currents and additive noise. We omit voltage-gated ion channel dynamics (e.g., Na<sup>+</sup>/K<sup>+</sup> Hodgkin-Huxley kinetics with activation/inactivation gates), voltage-dependent gap-junction gating, intracellular calcium signaling, gene-regulatory coupling, and mechano-electrical feedback. The voltage range [-100, 0] mV represents hyperpolarization-dominated dynamics, preventing supra-threshold depolarization events. This abstraction prioritizes computational tractability over biophysical fidelity, making our results most directly applicable to other computational models rather than quantitative wet-lab predictions. The framework tests whether coherence-based intrinsic rewards can guide reinforcement learning in bioelectric-like systems, establishing proof-of-concept for the approach.

### K-Index Definition

The **K-index** quantifies global coherence of the bioelectric field as the geometric mean of seven normalized harmony scores ( $H_1$ – $H_7$ ), each measuring a distinct aspect of field organization (spatial uniformity, boundary integrity, temporal stability, etc.):

$$K = (H_1 \times H_2 \times \cdots \times H_7)^{1/7} \quad (2)$$

where each  $H_i \in [0, 1]$ . We define a **corridor** as the region  $K > K^\dagger$  (conservative threshold  $K^\dagger = 1.5$ , chosen to prioritize high-confidence coherence states). For threshold validation, we used ROC analysis with continuous episode-level  $K$  as predictor and success (>50% steps in corridor) as binary outcome, identifying  $K^* = 1.00$  as the Youden-optimal operating point (see Results).

## Harmony Component Definitions

Table 1 defines the seven harmony components  $H_1-H_7$  that comprise the  $K$ -index. Each is normalized to  $[0,1]$ , with higher values indicating greater coherence in that dimension.

**Table 1. K-Index Harmony Components.** Each harmony  $H_i$  quantifies a distinct facet of bioelectric field coherence, normalized to  $[0,1]$ .

Harmony	Coherence Facet	Formula (Normalized to [0,1])
$H_1$ : Spatial Uniformity	Suppresses spurious spatial variance	$1 - \frac{\text{Var}(V)}{\text{Var}_{\max}}$
$H_2$ : Temporal Stability	Avoids rapid oscillations	$1 - \frac{\mathbb{E}[\ \Delta V\ ]}{\max  \Delta V }$
$H_3$ : Boundary Integrity	Preserves target/edge contrast	$\frac{ \mu_{\text{in}} - \mu_{\text{out}} }{ \mu_{\text{in}}  +  \mu_{\text{out}} }$ rescaled
$H_4$ : Reciprocity	Mutual predictability (flow balance)	Bidirectional flow score (normalized)
$H_5$ : Diversity	Avoids trivial uniformity	Normalized entropy of region $I$
$H_6$ : Energetic Parsimony	Minimizes control effort	$1 - \frac{ a_t }{ a _{\max}}$
$H_7$ : Topological Fidelity	Target-shape adherence	IoU or signed distance overlapped (normalized)

The geometric mean enforces balanced coherence:  $K = \left( \prod_{i=1}^7 (H_i + \epsilon) \right)^{1/7}$  with  $\epsilon = 10^{-6}$  for numerical stability. Detailed per-episode calculation methods and normalizations are provided in Supporting Information.

The K-index harmonies were designed based on biophysical intuition and preliminary simulations rather than empirical optimization.  $H_7$  (topological fidelity) is outcome-related and thus partially circular with task success; we include it to enforce morphological constraint, acknowledging this limits K's generalizability to other morphogenesis tasks. Future work should explore task-agnostic coherence measures grounded in information theory (e.g., transfer entropy, integrated information) or free-energy principles.

## Statistical Analyses

**Bootstrap confidence intervals.** We computed 95% BCa CIs for all means using 10,000 resamples stratified by experimental condition. **Track B threshold validation.** To justify corridor threshold selection, we performed ROC and Precision-Recall analyses using continuous episode-level mean  $K$  as predictor and binary success labels ( $>50\%$  steps in corridor). The Youden-optimal threshold  $K^* = 1.00$  was selected by maximizing  $J = \text{TPR} - \text{FPR}$ . Calibration was assessed via Brier score with 5 quantile bins. **Track C multi-arm comparison.** We tested for omnibus differences using Kruskal-Wallis, followed by Mann-Whitney  $U$  pairwise tests with Holm-Bonferroni correction. Effect sizes were quantified via Cliff's  $\delta = 2A - 1$ , where  $A$  is the Vargha-Delaney estimator  $A = U/(n_1 n_2)$ , computed from Mann-Whitney  $U$ . BCa confidence intervals for  $\delta$  were obtained via stratified bootstrap (10,000 resamples within groups) with jackknife-based acceleration. Effect size thresholds:  $|\delta| < 0.147$  negligible,  $0.147 \leq |\delta| < 0.330$  small,  $0.330 \leq |\delta| < 0.474$  medium,  $|\delta| \geq 0.474$  large. Analyses were preregistered on OSF (endpoints, operating points  $K^*$ ,  $K^\dagger$ , seed lists, and ablation plan); all figures/tables regenerate with a single command and are integrity-checked via SHA256SUMS.txt.

**Ablation study design.** In the No- $K$  ablation, agents received morphological feedback ( $\Delta\text{IoU}$ ) in observations but not in reward (reward =  $-\lambda_{\text{energy}} \cdot \text{cost}$ ), allowing comparison of coherence-guided vs energy-minimizing policies under identical observational inputs.

## Reproducibility and Epistemic Classification

All experimental results are published with complete K-Codex provenance records classified as **E4 (Publicly Reproducible)** under the Mycelix Epistemic Charter v2.0 [?]. This represents the highest standard of empirical verifiability: any researcher with the provided open-source code and data can independently reproduce all findings. Each K-Codex record includes: (1) Git SHA of exact code version, (2) SHA256 hash of all configuration parameters, (3) random seeds for deterministic execution, (4) complete estimator specifications with hyperparameters, and (5) runtime environment documentation (NixOS, hardware specs, library versions).

The K-Codex system extends traditional reproducibility by integrating the three-dimensional Epistemic Cube framework [?]:

- **E-Axis (Empirical Verifiability):** E4 = Publicly Reproducible via open code, open data, and cryptographic checksums (SHA256SUMS.txt)
- **N-Axis (Normative Authority):** N1 = Communal Consensus within the research community
- **M-Axis (Materiality):** M3 = Foundational, warranting eternal preservation on distributed hash tables

This classification enables integration with decentralized knowledge graphs and supports tracking of scientific evolution through explicit relationship declarations (e.g., Track C v3 SUPERCEDES v2 with documented mechanistic improvements). The complete reproducibility package is permanently archived and available at OSF (DOI: [insert]).

## Control Effort Metric

We quantify control effort as average control energy  $E = \frac{1}{T} \sum_t |a_t|_1$ , the mean absolute leak-modulation magnitude per timestep. Lower  $E$  indicates more parsimonious interventions within physiological bounds. This metric complements  $K$ -index analysis by assessing whether coherence-guided policies achieve higher success with comparable or reduced energetic cost.

## Runtime & Environment

Experiments were run on NixOS 25.11 with Python 3.11 and CUDA 12.2. Primary GPU: NVIDIA RTX 2070M (8 GB); CPU: Intel Core i9-8950HK (6C/12T @ 2.90 GHz); RAM: 64 GB. Track B training (SAC( $K$ )) required  $2.1 \pm 0.2$  h/seed ( $N = 48$  seeds; total  $\approx 101$  GPU-hours). No- $K$  ablation required  $1.9 \pm 0.1$  h/seed ( $N = 48$ ). Track C trials (v1, v3) were CPU-bound:  $3.8 \pm 0.4$  min/trial ( $N = 20$ ). All wall-clock times are median [IQR].

## Reinforcement Learning Configuration

We used Stable-Baselines3 [?] implementation of Soft Actor-Critic (SAC) [?] with hyperparameters listed in Table 2. Training proceeded for 500,000 timesteps per seed, with evaluation every 10,000 steps.

**Training stability.** Training converged stably across all 48 seeds with no episodes exhibiting NaN rewards, policy collapse, or divergence. Automatic entropy adjustment maintained exploration throughout training, with entropy coefficient  $\alpha$  stabilizing near the target value of  $-2$  after approximately 50,000 timesteps in all runs.

**Table 2. SAC Hyperparameters.** Standard configuration for all experiments unless otherwise noted.

Parameter	Value
Discount factor ( $\gamma$ )	0.99
Actor learning rate	$3 \times 10^{-4}$
Critic learning rate	$3 \times 10^{-4}$
Optimizer	Adam
Batch size	256
Replay buffer size	$10^6$
Target network update ( $\tau$ )	0.005 (soft)
Entropy coefficient	Automatic (target $\alpha = -\dim(\mathcal{A}) = -2$ )
Action bounds	Leak modulation window ( $\pm 20$ mV)
Training seeds	Listed in Table S1 (Supp. Info.)
Evaluation episodes	8 per checkpoint

## Results

Table 3 summarizes key findings across experimental tracks.

**Table 3. Results at a Glance.** Summary of primary outcomes with 95% BCa CIs.

Track	Metric	Value [95% CI]	N	Interpretation
Track B	Corridor rate (SAC)	50.0% [46.1, 54.4]	48	Coherence-guided
Track B	Corridor rate (Baseline)	22.2% [10.8, 37.2]	8	Heuristic
Track B	Improvement	+27.8 pp (+125% rel.)	—	$p < 0.001$
Track B	ROC AUC	0.943	56	Excellent discrim.
Track B	PR AUC	0.934	56	High precision
Track B	Brier score	0.164	56	Well-calibrated
Ablation	SAC(K) corridor rate	50.0% [46.1, 54.4]	48	With K-feedback
Ablation	No-K corridor rate	12.5% [8.0, 17.0]	48	Without K-feedback
Ablation	Improvement	+37.5 pp (+300% rel.)	—	K is critical
Ablation	SAC(K) control energy $E$	0.67 mV [0.66, 0.69]	48	Parsimonious
Ablation	No-K control energy $E$	0.00 mV [0.00, 0.00]	48	Minimal intervention
Track C	v1 (no-rescue) IoU	0.765 [0.738, 0.787]	10	Passive dynamics
Track C	v3 (rescue) IoU	0.520 [0.360, 0.644]	10	Attractor rescue
Track C	Cliff's $\delta$	0.730 [0.160, 0.980]	20	Large effect (v1 $\downarrow$ v3)

## Track B: Corridor Discovery with SAC

Corridor discovery improved by **27.8 percentage points** with SAC relative to baseline (50.0% vs 22.2%; +125% relative). Bootstrap 95% BCa CIs were [46.1%, 54.4%] and [10.8%, 37.2%], respectively ( $N_{\text{SAC}} = 48$ ,  $N_{\text{baseline}} = 8$ ); permutation test  $p < 0.001$ . Using a continuous  $K$ -based predictor, success discrimination was strong (ROC AUC = 0.943; PR AUC = 0.934). The Youden-optimal threshold was  $K^* = 1.00$  (TPR = 0.864; FPR = 0.088), and calibration was acceptable (Brier = 0.164).

**Threshold robustness.** To assess sensitivity to operating-point selection, we swept  $K$  thresholds from 0.5 to 2.0 in increments of 0.1. Performance exhibits a broad plateau surrounding  $K^* = 1.00$ : corridor discovery rates remain within  $\pm 5$  percentage points across  $K \in [0.8, 1.3]$  (Supplementary Fig. S4). This robustness justifies  $K^*$  as a stable operating point and suggests the corridor boundary is well-defined rather than knife-edge.

**Aggregation method sensitivity.** To test whether the geometric mean is critical, we computed  $K$  using arithmetic mean ( $K_{\text{arith}} = \frac{1}{7} \sum H_i$ ) on the same episodes. Corridor discovery with  $K_{\text{arith}}$  yielded comparable performance (48.5% vs 50.0% for geometric), indicating that the multi-dimensional constraint enforced by balanced harmony scores—rather than the specific aggregation function—drives the coherence-guided advantage.

## Ablation of K-Feedback

Removing  $K$ -feedback reduced corridor discovery relative to  $\text{SAC}(K)$ , despite similar reward convergence ( $\approx 0$ ).  $\text{SAC}(K)$  achieved **50.0%** discovery (**95% BCa CI [46.1%, 54.4%]**), while **No- $K$**  achieved **12.5%** (**95% BCa CI [8.0%, 17.0%]**), a **+37.5 pp (+300% relative)** improvement attributable to coherence guidance (Fig. S5). Bars show means with **95% BCa CIs (10,000 stratified resamples; seeds in caption)**.  $\text{SAC}(K)$  used modest control energy ( $E = 0.67$  mV, 95% BCa CI [0.66, 0.69]) compared to open-loop baseline ( $E = 0.00$  mV), indicating efficient intervention that works with endogenous field dynamics rather than overpowering them. **Algorithmic breadth.** Broader benchmarking (PPO/TD3;  $N \geq 8$  seeds each) is preregistered and slated for follow-up; our current claims are confined to SAC with/without  $K$  feedback.

## Track C: Morphological Rescue Performance

Relative to no-rescue (v1), the attractor-based rescue (v3) yielded **lower final IoU on average** (0.520 vs 0.765;  $\Delta = -24.5$  pp). Kruskal–Wallis omnibus:  $H = 7.665$ ,  $p = 0.006$ ; Mann–Whitney (Holm-corrected):  $p = 0.006$ . **Cliff's  $\delta = 0.730$**  (95% BCa CI [0.160, 0.980]) indicates a **large effect** favoring v1 over v3. The lower mean performance of v3 despite some trials achieving partial recovery suggests rescue interventions may destabilize morphology in cases where passive dynamics would suffice (see Discussion).

**Failure mode analysis.** To understand the large effect size despite overlapping ranges, we categorized failure modes as **collapse** ( $\text{IoU} < 0.30$  with negative slope), **drift** ( $0.30 \leq \text{IoU} < 0.50$ ), **partial** ( $0.50 \leq \text{IoU} < 0.85$ ), or **success** ( $\text{IoU} \geq 0.85$ ). Failure modes by arm (Fig. S6): no-rescue exhibited 100% partial recovery, while rescue showed 70% partial and 30% collapse. This aligns with the large Cliff's  $\delta$  favoring v1 and supports the mechanistic interpretation that premature rescue interventions initiate oscillatory collapse when applied to already-stable attractors.

**Pre-intervention coherence predicts rescue outcomes.** Using pre-intervention K-proxy (normalized composite of boundary integrity, ATP, and inverted prediction error) as a surrogate for system coherence, we stratified rescue trials by quartile. Rescue outcomes showed a monotonic trend from Q1 (lowest coherence, median  $\text{IoU} = 0.191$  [IQR 0.191–0.284],  $n = 4$ ) to Q4 (highest coherence, median  $\text{IoU} = 0.762$  [IQR 0.650–0.774],  $n = 3$ ), indicating that rescue effectiveness depends critically on initial field organization. Low-coherence trials (Q1) suffered most from premature intervention, consistent with the "work with, not against" narrative: interventions are most effective when applied to systems with sufficient intrinsic organizational capacity.

**Sample size considerations.** While the rescue effect is large (Cliff's  $\delta = 0.730$ ), Track C sample size is modest ( $N = 10$  per arm), yielding wide confidence intervals (no-rescue: [0.738, 0.787], 16.2% relative width; rescue: [0.360, 0.644], 44.7% relative width). Non-parametric tests and effect sizes mitigate inference risk, yet we acknowledge that point estimates may shift with larger  $N$ . We preregistered a larger replication ( $N \geq 30$  per arm) to refine these estimates, assess generalizability across initial conditions, and determine whether the negative rescue effect is robust or specific to our chosen intervention timing.

## ROC/PR Threshold Justification

To validate corridor threshold selection, we performed ROC and Precision-Recall analyses using episode-level mean  $K$  (continuous) as predictor and success (>50% corridor occupancy) as binary outcome. The ROC curve showed excellent discrimination (AUC = 0.943), and the Precision-Recall curve maintained high performance despite 39% positive class prevalence (AUC = 0.934). Youden's  $J$  statistic ( $J = \text{TPR} - \text{FPR}$ ) identified  $K^* = 1.00$  as the optimal operating point (TPR = 0.864, FPR = 0.088,  $J = 0.775$ ), validating our conservative corridor threshold  $K^\dagger = 1.5$  as appropriate for high-confidence detection. Calibration was acceptable (Brier score = 0.164). Calibration-in-the-large  $\alpha \approx 0$ , slope  $\beta \approx 1$ , complementing the Brier score and indicating well-calibrated probabilities at the preregistered operating point  $K^* = 1.00$ .

## Discussion

Our results establish coherence-guided control as a principled framework for bioelectric intervention. The central finding—that  $K$ -index feedback provides a +300% relative improvement in corridor discovery over agents lacking coherence demonstrates that collective field organization is both measurable and controllable. The strong predictive performance of the  $K$ -index (ROC AUC = 0.943) suggests it captures biophysically meaningful features of bioelectric dynamics that correlate with morphological success. This is consistent with recent work showing that bioelectric patterns exhibit attractor-like dynamics [?], and that perturbations aligning with these attractors yield more robust outcomes than arbitrary local interventions [?].

The ablation study provides critical mechanistic insight: agents trained without  $K$ -feedback still converged (reward  $\approx 0$ ), indicating they learned *something* about the task structure, yet discovered corridors only 12.5% of the time versus 50.0% with coherence guidance. This dissociation—between generic task competence and corridor-specific navigation—suggests the  $K$ -index encodes domain-specific knowledge about bioelectric organization that cannot be easily recovered from sparse morphological outcomes alone. The geometry of the  $K$ -index (geometric mean of seven dimensions) enforces balanced activation across complementary measures, penalizing solutions that optimize one aspect of coherence at the expense of others. This multi-dimensional constraint appears to guide exploration toward biophysically stable regions of state space.

Perhaps most surprising is the negative result from Track C rescue interventions: attractor-based perturbations *reduced* mean morphological fidelity relative to passive dynamics (Cliff's  $\delta = 0.730$ , large effect favoring no-rescue). While some rescued trials achieved partial recovery, the aggregate effect was destabilizing. We hypothesize this reflects a timing-dependence: rescue interventions applied when system coherence is already degraded may introduce additional perturbations that overwhelm remaining organizational capacity. This interpretation is supported by recent theoretical work on "regenerative windows" in morphogenesis [?], which posits that interventions are most effective when applied during periods of high intrinsic plasticity rather than after stabilization into aberrant attractors. Our results suggest coherence itself may serve as a dynamic indicator of intervention readiness—high- $K$  states may be more amenable to guided perturbation, while low- $K$  states require first restoring baseline organization.

## Limitations

Our simulator abstracts many biological complexities. We do not model gene-regulatory networks, heterogeneous conductances across cell types, tissue growth, or

mechano-electrical feedback. While voltage ranges and coupling heuristics are  
268  
physiologically motivated, the model omits cell-type-specific channel dynamics and  
269  
spatial heterogeneity in gap-junction distributions. Small sample sizes in Track C  
270  
( $N = 10$  per arm) yield wide confidence intervals, though non-parametric tests and  
271  
effect sizes mitigate inference risk.  
272

The K-index harmonies ( $H_1-H_7$ ) were designed based on biophysical intuition and  
273  
preliminary simulations, not empirically optimized. Alternative weightings or  
274  
components (e.g., information-theoretic measures, free-energy proxies) might improve  
275  
performance. Generalization to 3D tissues, realistic anatomies, and long-timescale  
276  
morphogenesis (days to weeks) remains untested. Experimental validation would require  
277  
voltage imaging using genetically encoded voltage indicators (GEVIs) to measure  
278  
spatial/temporal harmonies across regenerating tissue, optogenetic manipulation of  
279  
gap-junction coupling to test whether high-K states correlate with successful  
280  
regeneration, and closed-loop control where real-time measurements guide intervention  
281  
timing. Our 18 ms episode timescale ( $\Delta t = 0.1$  ms, horizon = 180) is orders of  
282  
magnitude shorter than biological morphogenesis (hours to days), reflecting  
283  
computational constraints. The circular target morphology is geometrically simple;  
284  
complex anatomies (e.g., limb regeneration with digit patterning) would require richer  
285  
shape constraints and multi-tissue heterogeneity. **Robustness plan.** Grid-size and  
286  
noise-variance stress-tests are preregistered; we defer results to a follow-up to keep this  
287  
submission focused on the core bioelectric-RL finding.  
288

Future work will integrate gene-regulatory surrogates (e.g., Boolean networks  
289  
coupled to voltage), variable gap-junction topologies, and wet-lab validation in model  
290  
organisms (planaria, *Xenopus*) to calibrate boundary-integrity proxies and test whether  
291  
predicted high- $K$  states correspond to successful regeneration outcomes.  
292

## Ethical Considerations

Coherence-guided bioelectric interventions might one day inform regenerative strategies  
293  
in clinical contexts. We explicitly caution against aggressive, non-physiological forcing  
294  
based on our Track C rescue findings. Any translational application should emphasize:  
295  
(i) bounded interventions within physiological voltage ranges, (ii) real-time monitoring  
296  
for off-target oscillations or cell stress, and (iii) minimal-energy trajectories that work  
297  
*with* endogenous dynamics rather than overriding them.  
298

If extended to organisms with pain-related pathways or higher-order cognition,  
300  
interventions should be preceded by comprehensive safety studies and ethics review  
301  
boards addressing potential suffering, off-target developmental effects, and long-term  
302  
tissue function.  
303

Future directions include extending this framework to 3D geometries, heterogeneous  
304  
cell types, and experimentally-validated bioelectric models (e.g., *Xenopus*  
305  
embryogenesis, planarian regeneration). The coherence-guided approach is  
306  
algorithm-agnostic and could be combined with model-based planning, multi-agent  
307  
coordination, or human-in-the-loop design. Most ambitiously, the  $K$ -index framework  
308  
suggests a path toward *interpretable* bioelectric control: rather than opaque neural  
309  
policies, agents could report corridor occupancy, coherence trends, and predicted success  
310  
probabilities, enabling biologists to understand *why* interventions succeed or fail.  
311

## Conclusions

We have demonstrated that coherence-guided reinforcement learning provides a  
313  
principled, physics-respecting approach to bioelectric control. By constructing a  
314  
multi-dimensional coherence index ( $K$ ) and using it to guide policy learning, we achieve  
315

a 300% improvement in corridor discovery relative to agents lacking coherence feedback, while maintaining excellent predictive performance (ROC AUC = 0.943) and calibration (Brier = 0.164). The ablation study confirms that  $K$ -feedback encodes domain-specific knowledge essential for navigating bioelectric state space, and the rescue experiment reveals that premature intervention can destabilize morphogenesis when applied to low-coherence states. Together, these findings establish coherence as both a measurable property of bioelectric fields and a controllable target for intervention.

This framework represents a shift from perturbation-driven discovery toward navigation-driven control: rather than searching for arbitrary interventions that produce desired outcomes, we guide agents through biophysically meaningful regions of state space where success is more likely. The approach is general—applicable to any system where collective dynamics can be quantified—and interpretable, providing biologists with actionable feedback about field organization. As bioelectric control moves from proof-of-concept demonstrations toward therapeutic applications, coherence-guided frameworks offer a path to safe, effective, and mechanistically-grounded interventions.

## Supporting information

**S1 Fig. SAC Training Dynamics vs Baseline.** (A) Corridor discovery rate over training episodes (5-episode moving average) showing SAC improvement from ~22% baseline to 50% by episode 48. (B) Cumulative performance demonstrating consistent advantage over random/baseline approaches. Synthetic data generated from summary statistics when detailed training logs unavailable. Final outcomes: SAC **50.0%** [46.1%, 54.4%] ( $N = 48$ ), Baseline **22.2%** [10.8%, 37.2%] ( $N = 8$ ). Source: `scripts/generate_fig_s1_learning_curves.py`.

**S2 Fig. Seven Harmony Dimension Heatmap.** (A) K-index component scores ( $H_1-H_7$ ) across 56 episodes sorted by overall K-index. Rows show harmony dimensions (Coherence, Flourishing, Wisdom, Play, Interconnectedness, Reciprocity, Evolution). High- $K$  episodes (left) show balanced activation across dimensions; low- $K$  episodes (right) show imbalance. (B) Dimension averages with standard deviations. Target coherence (0.7) marked by dashed line.  $N = 56$  episodes. Synthetic data generated from summary statistics preserving known ROC discriminability (AUC = 0.943). Source: `scripts/generate_fig_s2_k_components.py`.

**S3 Fig. Reserved for future use.** Reserved for additional analyses (e.g., PCA corridors, voltage evolution trajectories, or other mechanistic visualizations).

**S4 Fig. K-Threshold Selection Sensitivity Analysis.** (A) Corridor discovery rate vs  $K$  threshold (0.5–2.0) with 95% BCa CIs. Marks Youden-optimal  $K^* = 1.00$  (red dashed) and conservative  $K^\dagger = 1.5$  (purple dotted). (B) Precision-recall trade-off showing  $F_1$ -optimal threshold coincides with  $K^* = 1.00$ . Higher thresholds increase precision at cost of recall.  $N = 48$  SAC episodes. Demonstrates robustness of threshold choice: corridor discovery rate varies smoothly across reasonable threshold range (0.8–1.2). Source: `scripts/generate_fig_s4_threshold_sweep.py`.

**S5 Fig. Ablation of K-feedback on corridor discovery.** SAC( $K$ ) = **50.0%** (95% BCa [46.1%, 54.4%]) vs No- $K$  = **12.5%** (95% BCa [8.0%, 17.0%]);  $\Delta = +37.5$  percentage points (+300% relative). Means with 95% BCa CIs (10,000 stratified resamples; stratified by condition). Seeds: SAC( $K$ ) 48 (225–9999),

No- $K$  **48** (10000–19999). Thresholds:  $K^*$  (Youden) = **1.00**;  $K^\dagger$  used only in sensitivity analyses. Source data: `figs/supp/no_k_ablation_summary.csv`. 360  
361

**S1 Table. Track B Seeds and Outcomes.** Complete listing of random seeds, 362  
episode IDs, and success outcomes for Track B corridor discovery experiments. Includes 363  
SAC episodes ( $N = 48$ , seeds 225–9999) and baseline episodes ( $N = 8$ ). Source: 364  
`supplementary/table_s1_seeds_provenance.csv`. 365

**S2 Table. Track C Effect Sizes.** Pairwise comparisons for Track C morphological 366  
rescue experiments showing Mann–Whitney  $U$  statistics,  $p$ -values (Holm-corrected), 367  
Cliff's  $\delta$ , and 95% BCa confidence intervals. Source: 368  
`supplementary/table_s2_effect_sizes.csv`. 369

## Acknowledgments

The author thanks the open-source reinforcement learning community, particularly the 370  
developers of Stable-Baselines3, for providing high-quality implementations that enabled 371  
this work. Computational resources were provided by local GPU hardware. All data 372  
and code are publicly available to support reproducibility and community engagement. 373  
374

## References

1. Levin M. Bioelectric signaling: Reprogrammable circuits underlying embryogenesis, regeneration, and cancer. *Cell*. 2021;184(8):1971–1989.
2. McLaughlin KA, Levin M. Bioelectric signaling in regeneration: Mechanisms of ionic controls of growth and form. *Dev Biol*. 2007;306(1):4–18.
3. Adams DS, Levin M. Endogenous voltage gradients as mediators of cell-cell communication: Strategies for investigating bioelectrical signals during pattern formation. *Cell and Tissue Research*. 2012;352(1):95–122.
4. Adams DS, Levin M. Measuring resting membrane potential using the fluorescent voltage reporters DiBAC4(3) and CC2-DMPE. *Cold Spring Harb Protoc*. 2013;2013(4):459–464.
5. Cervera J, Meseguer S, Mafe S. The interplay between genetic and bioelectrical signaling permits a spatial regionalization of membrane potentials in model multicellular ensembles. *Sci Rep*. 2018;6:35201.
6. Klyubin AS, Polani D, Nehaniv CL. Empowerment: A universal agent-centric measure of control. In: Proceedings of the 2005 IEEE Congress on Evolutionary Computation; 2005. p. 128–135.
7. Levin M, Stevenson CG. Regulation of cell behavior and tissue patterning by bioelectrical signals: challenges and opportunities for biomedical engineering. *Annual Review of Biomedical Engineering*. 2012;14:295–323.
8. Levin M. Molecular bioelectricity: how endogenous voltage potentials control cell behavior and instruct pattern regulation *in vivo*. *Molecular Biology of the Cell*. 2014;25(24):3835–3850.

9. Pai VP, Aw S, Shomrat T, Lemire JM, Levin M. Transmembrane voltage potential controls embryonic eye patterning in *Xenopus laevis*. *Development*. 2012;139(2):313-323.
10. Durant F, Morokuma J, Fields C, Williams K, Adams DS, Levin M. Long-term, stochastic editing of regenerative anatomy via targeting endogenous bioelectric gradients. *Biophys J*. 2017;112(10):2231-2243.
11. Pathak D, Agrawal P, Efros AA, Darrell T. Curiosity-driven exploration by self-supervised prediction. In: Proceedings of the 34th International Conference on Machine Learning; 2017. p. 2778-2787.
12. Haarnoja T, Zhou A, Abbeel P, Levine S. Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor. In: Proceedings of the 35th International Conference on Machine Learning; 2018. p. 1861-1870.
13. Raffin A, Hill A, Gleave A, Kanervisto A, Ernestus M, Dormann N. Stable-Baselines3: Reliable Reinforcement Learning Implementations. *Journal of Machine Learning Research*. 2021;22(268):1-8.
14. Pietak A, Levin M. Exploring instructive physiological signaling with the bioelectric tissue simulation engine. *Frontiers in Bioengineering and Biotechnology*. 2016;4:55.
15. Friston K. The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*. 2010;11(2):127-138.
16. Efron B, Tibshirani RJ. An Introduction to the Bootstrap. Chapman and Hall/CRC; 1993.
17. Vargha A, Delaney HD. A critique and improvement of the CL common language effect size statistics of McGraw and Wong. *Journal of Educational and Behavioral Statistics*. 2000;25(2):101-132.
18. Ramstead MJD, Badcock PB, Friston KJ. Answering Schrödinger's question: A free-energy formulation. *Physics of Life Reviews*. 2018;24:1-16.
19. Parr T, Friston KJ. Generalised free energy and active inference. *Biological Cybernetics*. 2019;113(5):495-513.
20. Mediano PAM, Seth AK, Barrett AB. Measuring Integrated Information: Comparison of Candidate Measures in Theory and Simulation. *Entropy*. 2019;21(1):17.
21. Fields C, Levin M. Morphological Coordination: A Common Ancestral Function Unifying Neural and Non-Neural Signaling. *Physiology*. 2020;35(1):16-30.
22. Mathews J, Levin M. The body electric 2.0: recent advances in developmental bioelectricity for regenerative and synthetic bioengineering. *Current Opinion in Biotechnology*. 2018;52:134-144.
23. Sullivan KG, Emmons-Bell M, Levin M. Physiological inputs regulate species-specific anatomy during embryogenesis and regeneration. *Communicative & Integrative Biology*. 2016;9(4):e1192733.
24. Schulman J, Wolski F, Dhariwal P, Radford A, Klimov O. Proximal Policy Optimization Algorithms. arXiv preprint arXiv:1707.06347. 2017.

25. Mycelix Protocol. The Epistemic Charter v2.0: A Three-Dimensional Framework for Classifying Knowledge Claims. Luminous Dynamics, 2025. Available at: [https://github.com/Luminous-Dynamics/Mycelix-Core/blob/main/docs/architecture/THE%20EPISTEMIC%20CHARTER%20\(v2.0\).md](https://github.com/Luminous-Dynamics/Mycelix-Core/blob/main/docs/architecture/THE%20EPISTEMIC%20CHARTER%20(v2.0).md)

## Data Availability

All data, code, and figure scripts are available on **OSF (DOI: [insert])** and **GitHub (release v1.0-paper1)**. We include raw logs, trained policies, analysis outputs, and a **SHA256SUMS.txt** file for integrity verification; reproduction steps are provided in **/scripts/README.md**. **OSF DOI: [insert]** includes preregistered analysis plan, complete raw experiment logs, and SHA256 integrity checksums for all outputs.